

AI POWERED DEEPPFAKE FACE DETECTION

¹Manjunath V, ² Ganesh

¹ Assistant Professor, ² Student

¹ Department of Computer Applications,

¹ Government Degree College, Yadgir

Adikavi Maharshi Valmiki University, Raichur, India

Abstract— Deepfake technology, which uses artificial intelligence (AI) to create hyper-realistic images and videos by manipulating faces, has become increasingly sophisticated. While this technology holds promise for entertainment and creative purposes, it also presents serious threats, such as the spread of misinformation, identity theft, and privacy violations. Detecting deepfake content has thus become crucial in maintaining trust in digital media.

The rapid advancement of deepfake technology has raised significant concerns regarding the authenticity of digital media, particularly in the context of face manipulation in videos and images. Deepfakes, which employ artificial intelligence (AI) to superimpose faces onto videos or images, have the potential to cause significant harm by spreading misinformation and undermining trust in media

Index Terms— CNN, ResNet-50, DeepFake Detection, FaceForensics++, DeepFakeTIMIT Dataset, Transfer Learning, Pretrained Models, Image Manipulation, Real-Time Detection, TensorFlow, PyTorch, Binary Classification, Pixel-level Anomaly Detection, Data Augmentation, Model Evaluation, Precision

I. INTRODUCTION

The integration of computer vision with natural language processing has enabled machines to interpret and describe visual content effectively. Image captioning, the process of generating textual descriptions for images, is an interdisciplinary task involving both visual understanding and language generation. When paired with voice output, the solution becomes more inclusive and practical, especially for accessibility-focused use cases.

This research proposes a deep learning-based image captioning system using the BLIP (Bootstrapping Language-Image Pre-training) model, which has demonstrated state-of-the-art performance in caption generation. Additionally, the integration of text-to-speech technology bridges the gap between visual recognition and auditory response.

Ref Fig.1

II. MATERIALS AND METHODS

For training and evaluation, a custom dataset was curated comprising high-resolution facial images categorized into Real and Deepfake classes.

- Each image underwent standard preprocessing steps, including resizing to 224×224 pixels,
- Normalization using EfficientNetB0's `preprocess_input()`, and color space conversion from RGB to BGR using OpenCV utilities.
- Data augmentation strategies such as rotation, flipping, and zooming were implicitly handled via ImageDataGenerator to enhance model generalization.

III. Model Architecture

- The backbone of the detection framework is the EfficientNetB0 model, pre-trained on ImageNet and fine-tuned for binary classification. To adapt it for the deepfake detection task:
- The base convolutional layers were frozen to retain learned features.
- A custom classification head was appended, consisting of a Global Average Pooling 2D layer, a fully connected Dense layer with ReLU activation, a Dropout layer (rate = 0.5) to mitigate overfitting, and a final sigmoid-activated output neuron to predict the authenticity score.
- Optimization was performed using the Adam optimizer with a learning rate of 0.0001, and binary crossentropy was employed as the loss function.

IV. Training Strategy

- The model was trained for 8 epochs on mini-batches of size 32, leveraging both the training and validation splits dynamically generated via ImageDataGenerator. Real-time performance monitoring was conducted by tracking both training and validation metrics across epochs.
- The training pipeline was orchestrated in TensorFlow 2.x environment with GPU acceleration enabled where available.

V. Application Deployment

- For real-world usability, an interactive frontend was built using Streamlit. The deployed application allows users to upload an image, after which the system:
- Applies the identical preprocessing pipeline as during training.
- Passes the image through the trained EfficientNetB0 model.
- Outputs a prediction label ("Real" or "Fake") with associated confidence.
- The lightweight, responsive interface ensures cross-platform compatibility without the need for additional installations.

VI. Tools, Libraries, and Frameworks

- The following tools and libraries were utilized to implement, train, and deploy the solution:
- TensorFlow 2.11 – for deep learning model training and inference.
- Keras API – for model definition and transfer learning integration.
- OpenCV (cv2) – for advanced image manipulation and preprocessing.
- NumPy – for numerical data handling.
- Streamlit – for rapid deployment of the web-based inference system.
- Matplotlib – for visualization of training history and performance metrics.
- Pillow (PIL) – for efficient image loading and resizing.
- All experiments were conducted on a workstation configured with an Intel Core i7 processor, 16GB RAM, and an NVIDIA GTX 1660 GPU.




Fig.1

Results

The proposed deepfake detection system demonstrated strong performance on the curated dataset. After training for eight epochs, the model achieved a training accuracy of 97.3% and a validation accuracy of 95.6%, indicating high generalization capability without significant overfitting.

The loss curves and accuracy curves plotted across epochs exhibited stable convergence, with no abrupt spikes, suggesting effective learning dynamics.


On unseen test images, the model consistently distinguished between real and fake faces, delivering binary classification outputs with high confidence (>95%) in most cases.

The deployed Streamlit application successfully processed uploaded images within 1–2 seconds, offering real-time feedback to users.

Predictions were reliable even when tested on varied lighting conditions and image qualities, affirming the model's robustness and real-world applicability.

Sample Output:

- Input: Deepfake facial image.
- Prediction: "● Fake Image Detected" with 98.7% confidence.
- Inference Time: ~1.4 seconds.
- Overall, the system demonstrated high accuracy, rapid inference, and strong robustness against moderate variations in image quality.

Ref Fig.2**Fig.2****Discussion**

The advancement of synthetic media, particularly deepfake technologies, has posed significant ethical and security challenges in the digital landscape. In this work, we proposed a lightweight yet robust AI-powered framework for real-time deepfake face detection using EfficientNetB0 as the backbone architecture. The results achieved demonstrate promising capabilities, but several critical aspects deserve deeper discussion.

- Firstly, the utilization of **EfficientNetB0** enabled high accuracy while maintaining computational efficiency, making the system suitable for deployment even on moderately powered devices. Freezing the convolutional base and fine-tuning only the classification head prevented overfitting during the training phase, as evidenced by the smooth convergence of validation metrics.
- Secondly, the system's ability to generalize well on unseen data can be attributed to the rigorous pre-processing pipeline, balanced dataset preparation, and the application of real-world image augmentations during training. However, performance under extreme adversarial conditions (such as heavily compressed videos or obfuscated facial features) remains an open challenge and an avenue for future enhancement.
- Furthermore, the Streamlit-based frontend provided a rapid and intuitive user interface for end-users, bridging the gap between sophisticated backend models and practical usability. The response time of approximately **1–2 seconds per inference** is considered highly acceptable for real-world applications, particularly for mobile or edge devices.
- Despite these strengths, the system currently operates solely on static images, limiting its applicability to dynamic video content where temporal inconsistencies can serve as additional features for detection. Integration of temporal deepfake detection mechanisms, possibly using RNN or Transformer-based architectures, could be an effective direction for evolution.
- Additionally, as deepfake generation techniques continue to evolve (GANs, Diffusion Models, etc.), it becomes imperative to periodically retrain or fine-tune the model on newer datasets to maintain detection efficacy against sophisticated synthetic media.

In essence, the research validates that compact convolutional architectures, when coupled with strategic pre-processing and balanced datasets, can yield effective defenses against deepfake media. Continuous innovation, dataset enrichment, and adversarial robustness testing will be essential to stay ahead in the arms race against synthetic content manipulation.

Conclusion

In this study, we developed an AI-driven system for the detection of deepfake facial imagery by leveraging the efficiency of transfer learning with EfficientNetB0. The model demonstrated high classification accuracy, low latency, and strong generalization capabilities across a balanced dataset of real and synthetic images.

The combination of lightweight architecture, intelligent pre-processing, and a user-friendly Streamlit deployment proves that practical deepfake detection solutions can be both powerful and accessible. However, given the rapid evolution of deepfake generation techniques, continuous model updates and integration of advanced temporal analysis will be crucial for sustaining performance over time.

Ultimately, this work underscores the potential of deep learning-based solutions in safeguarding digital media authenticity and lays a solid foundation for future research towards more resilient, scalable, and real-time deepfake detection systems.

REFERENCES

- [1] Korshunov, P., & Marcel, S. (2018). *Deepfakes: a new threat to face recognition? Assessment and detection*. arXiv preprint arXiv:1812.08685.
- [2] Rossler, A., Cozzolino, D., Verdoliva, L., Riess, C., Thies, J., & Nießner, M. (2019). *FaceForensics++: Learning to detect manipulated facial images*. Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV).
- [3] Afchar, D., Nozick, V., Yamagishi, J., & Echizen, I. (2018). *MesoNet: A Compact Facial Video Forgery Detection Network*. Proceedings of the IEEE International Workshop on Information Forensics and Security (WIFS).
- [4] Nguyen, H. H., Yamagishi, J., & Echizen, I. (2019). *Capsule-forensics: Using capsule networks to detect forged images and videos*. ICASSP 2019 - IEEE International Conference on Acoustics, Speech and Signal Processing.
- [5] Tolosana, R., Vera-Rodriguez, R., Fierrez, J., Morales, A., & Ortega-Garcia, J. (2020). *Deepfakes and beyond: A survey of face manipulation and fake detection*. Information Fusion, 64, 131–148.
- [6] Chollet, F. (2017). *Xception: Deep learning with depthwise separable convolutions*. Proceedings of the IEEE conference on computer vision and pattern recognition (CVPR).